

# Cache Friendly Programming

**Oğuzhan KATLI**



<https://www.linkedin.com/in/ogzhnktl>

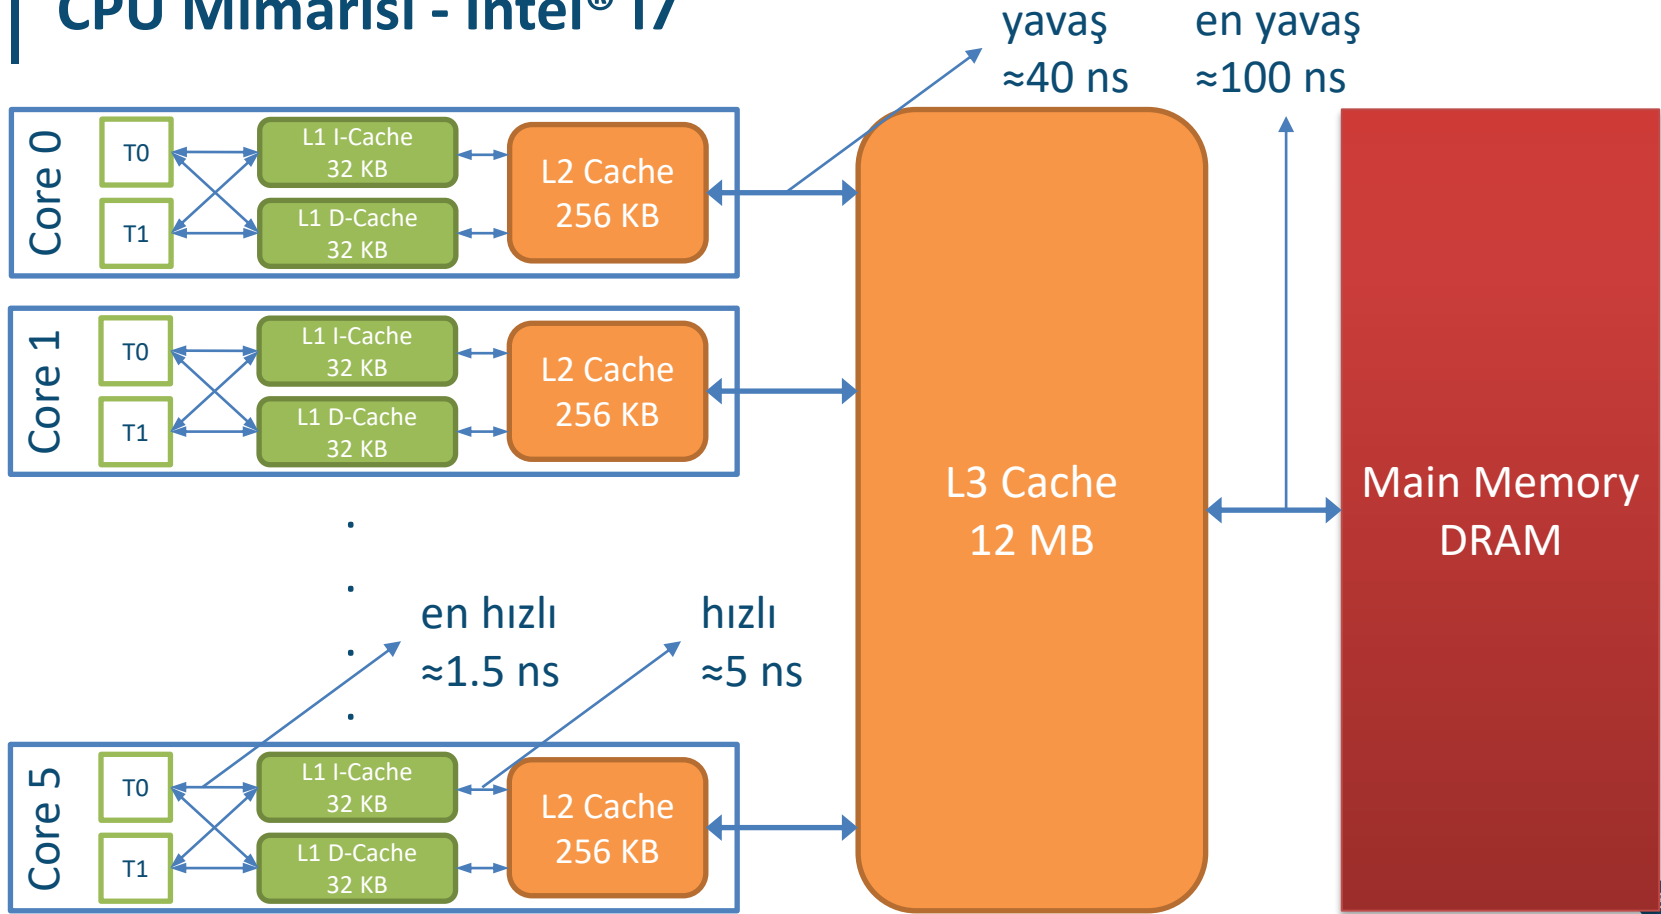


<https://github.com/nixiz>

# Benchmark Setup


- Intel® Core™ i7-9750H İşlemci (12M Önbellek, 4,50 GHz'e kadar)
  - 6 Core – 12 Threads
  - 32 KB L1 Cache
  - 256 KB L2, 12 MB L3 Cache
- Windows 10 Home
  - Visual Studio 2019 Version 16.7.2 – MSVC 19.27
- Ubuntu 20.04 LTS
  - GCC 9.3.0
  - Clang 10.0
- Google Benchmark: <https://github.com/google/benchmark>
- Benchmark Kodları: <https://gist.github.com/nixiz/c3bb8f6029b64f9281ac44cbc119209a>

# CPU Mimarisi - Intel® i7




# CPU Mimarisi – Cache Line

Okuma:



Yazma:



...

# Veri Erişim Hızları

## Latency Numbers Every Programmer Should Know

■ 1 ns

■ L1 cache reference: 0.5 ns

■ Branch mispredict: 5 ns

■ L2 cache reference: 7 ns

■ Mutex lock/unlock: 25 ns

■ = 100 ns

■ Main memory reference: 100 ns

■ = 1  $\mu$ s

■ Compress 1 KB with Zippy: 3  $\mu$ s

■ = 10  $\mu$ s

■ Send 1 KB over 1 Gbps network: 10  $\mu$ s

■ SSD random read (1 Gb/s SSD): 150  $\mu$ s

■ Read 1 MB sequentially from memory: 250  $\mu$ s

■ Round trip in same datacenter: 500  $\mu$ s

■ = 1 ns

■ Read 1 MB sequentially from SSD: 1 ns

■ Disk seek: 10 ms


■ Read 1 MB sequentially from disk: 20 ms

■ Packet roundtrip CA to Netherlands: 150 ns


Source: <https://gist.github.com/2841832>

# Row vs Column Matrix Traversal

# Row vs Column Major



Row Major




Column Major

# Row vs Column Major


```
template <typename T>
auto sumMatrix(const Matrix<T>& m, TraversalOrder order)
{
    unsigned long sum = 0;
    if (order == TraversalOrder::RowMajor)
    {
        for (auto r = 0; r < m.rows(); ++r)
            for (auto c = 0; c < m.columns(); ++c)
                sum += m[r][c];
    }
    else // TraversalOrder::ColumnMajor
    {
        for (auto c = 0; c < m.columns(); ++c)
            for (auto r = 0; r < m.rows(); ++r)
                sum += m[r][c];
    }
    return sum;
}
```




# Row vs Column Major Sonuç



# Row vs Column Major Sonuç



# Row vs Column Major Sonuç



# Branch Prediction

# Branch Prediction

v:	-1	-1	1	1	-1	1	-1	-1	1	-1	-1	1	1	-1	1
v:	1	1	1	1	1	1	1	-1	-1	-1	-1	-1	-1	-1	-1

```
std::vector<float> v(65536);  
std::generate(std::begin(v), std::end(v), [] {  
    return (rand() % 2) ? 1 : -1;  
});
```

```
std::sort(v.begin(), v.end());
```

```
return std::count_if(std::begin(v), std::end(v), [](float x) {  
    return x > 0;  
});
```


## Branch Prediction – Sequenced If - Else

v:	1	-1	1	-1	1	-1	1	-1	1	-1	1	-1	1	-1	1
----	---	----	---	----	---	----	---	----	---	----	---	----	---	----	---

```
std::vector<float> v(65536);
std::generate(std::begin(v), std::end(v), [n = 0]() mutable {
    return (++n % 2) ? 1 : -1;
});

return std::count_if(std::begin(v), std::end(v), [](float x) {
    return x > 0;
});
```

## Count If Sonuç



# Virtual Function Calls

```
struct base_price
{
    virtual ~price() {}
    virtual float getPrice() const noexcept { return 0.0; }
};

struct cheap : public base_price
{
    float getPrice() const noexcept override { return 2.0; }
};

struct expensive : public base_price
{
    float getPrice() const noexcept override { return 3.14159; }
};
```




# Örnek

```
std::vector<base_price*> pricelist;  
std::fill_n(std::back_inserter(pricelist), 10000, new base_price);  
std::fill_n(std::back_inserter(pricelist), 10000, new cheap);  
std::fill_n(std::back_inserter(pricelist), 10000, new expensive);
```

```
std::random_shuffle(pricelist.begin(), pricelist.end());
```

```
float sum = 0;  
for (auto *p : pricelist)  
{  
    sum += p->getPrice();  
}
```

# Virtual Call Sonuç



# False Sharing

# False Sharing

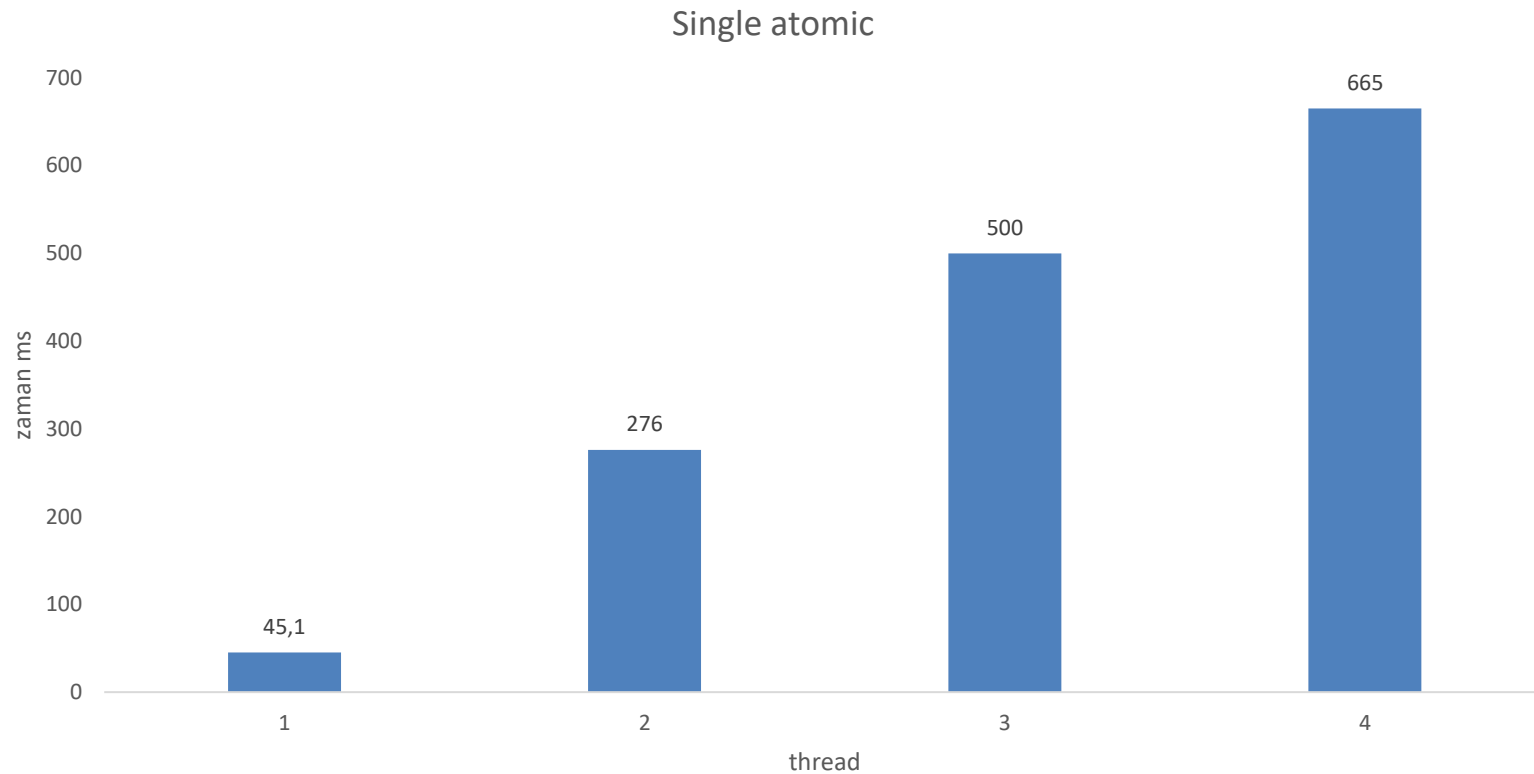
```
template <typename input_t>
void work(input_t& a)
{
    for (int i = 0; i < 10,000,000; ++i)
        a++;
}

int test_with_4_threads()
{
    std::atomic<int> a; a = 0;

    std::thread t1([&] { work(a); });
    std::thread t2([&] { work(a); });
    std::thread t3([&] { work(a); });
    std::thread t4([&] { work(a); });

    t1.join(); t2.join(); t3.join(); t4.join();
    return a;
}
```

# Sonuç



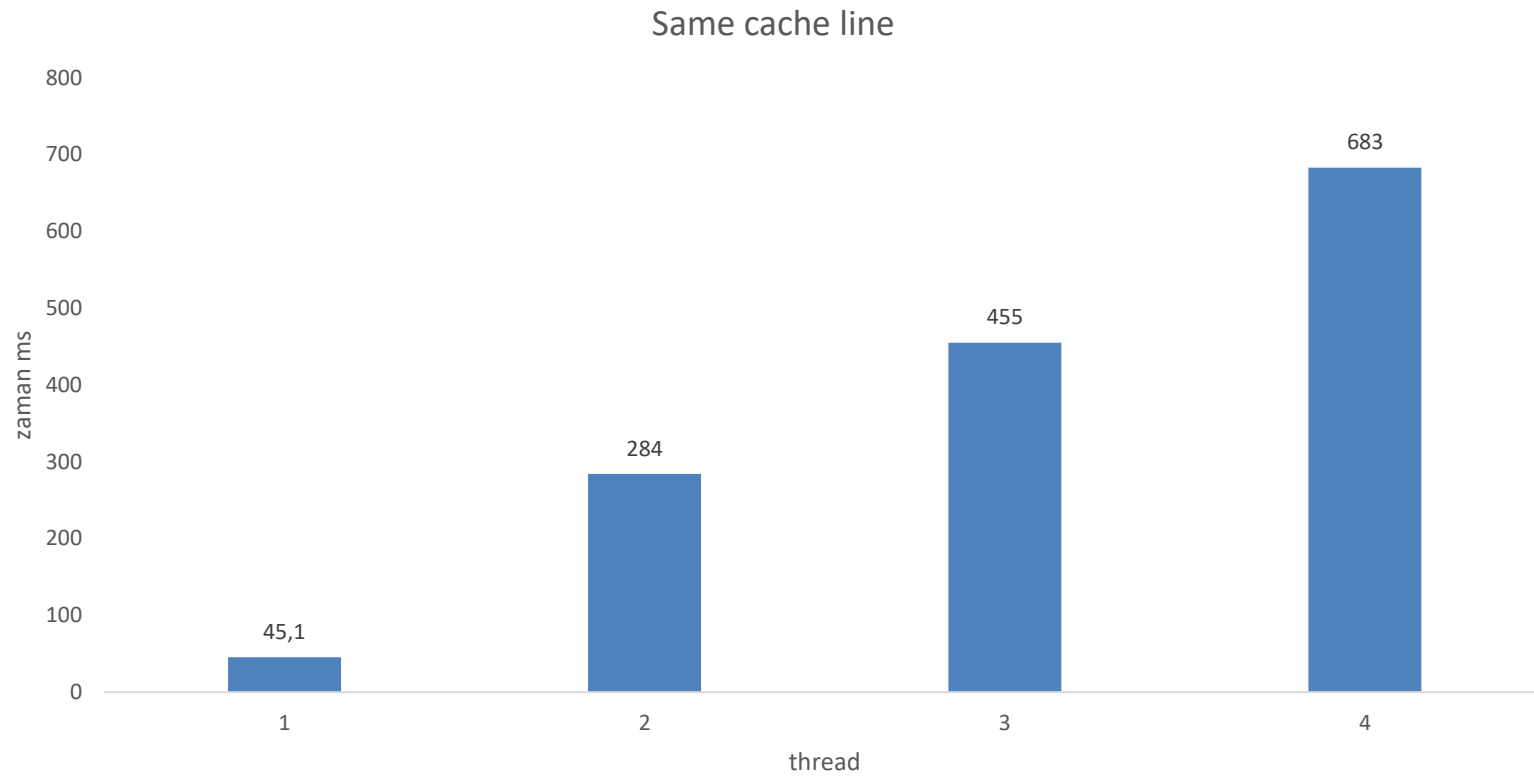
# Farklı Atomik Nesneler

```
int test()
{
    std::atomic<int> a; a = 0;
    std::atomic<int> b; b = 0;
    std::atomic<int> c; c = 0;
    std::atomic<int> d; d = 0;

    std::thread t1([&] { work(a); });
    std::thread t2([&] { work(b); });
    std::thread t3([&] { work(c); });
    std::thread t4([&] { work(d); });

    t1.join(); t2.join(); t3.join(); t4.join();
    return a + b + c + d;
}
```

# Sonuç



# Farklı Atomik Nesneler

```
int test()
{
    std::atomic<int> a; a = 0; // &a: 0x...b2f7c0
    std::atomic<int> b; b = 0; // &b: 0x...b2f7c4
    std::atomic<int> c; c = 0; // &c: 0x...b2f7c8
    std::atomic<int> d; d = 0; // &d: 0x...b2f7cc

    std::thread t1([&] { work(a); });
    std::thread t2([&] { work(b); });
    std::thread t3([&] { work(c); });
    std::thread t4([&] { work(d); });

    t1.join(); t2.join(); t3.join(); t4.join();
    return a + b + c + d;
}
```

Adres blokları aynı



# Çözüm

```
struct alignas(64) aligned_type
{
    std::atomic<int> val;
};
```

```
int test()
```

```
{
    aligned_type a; a.val = 0; // &a: 0x...4ff240
    aligned_type b; b.val = 0; // &b: 0x...4ff280
    aligned_type c; c.val = 0; // &c: 0x...4ff2c0
    aligned_type d; d.val = 0; // &d: 0x...4ff300
```


Adres blokları Farklı

```
    std::thread t1([&a] { work(a.val); });
    std::thread t2([&b] { work(b.val); });
    std::thread t3([&c] { work(c.val); });
    std::thread t4([&d] { work(d.val); });
```

```
    t1.join(); t2.join(); t3.join(); t4.join();
    return a.val + b.val + c.val + d.val;
```

```
}
```

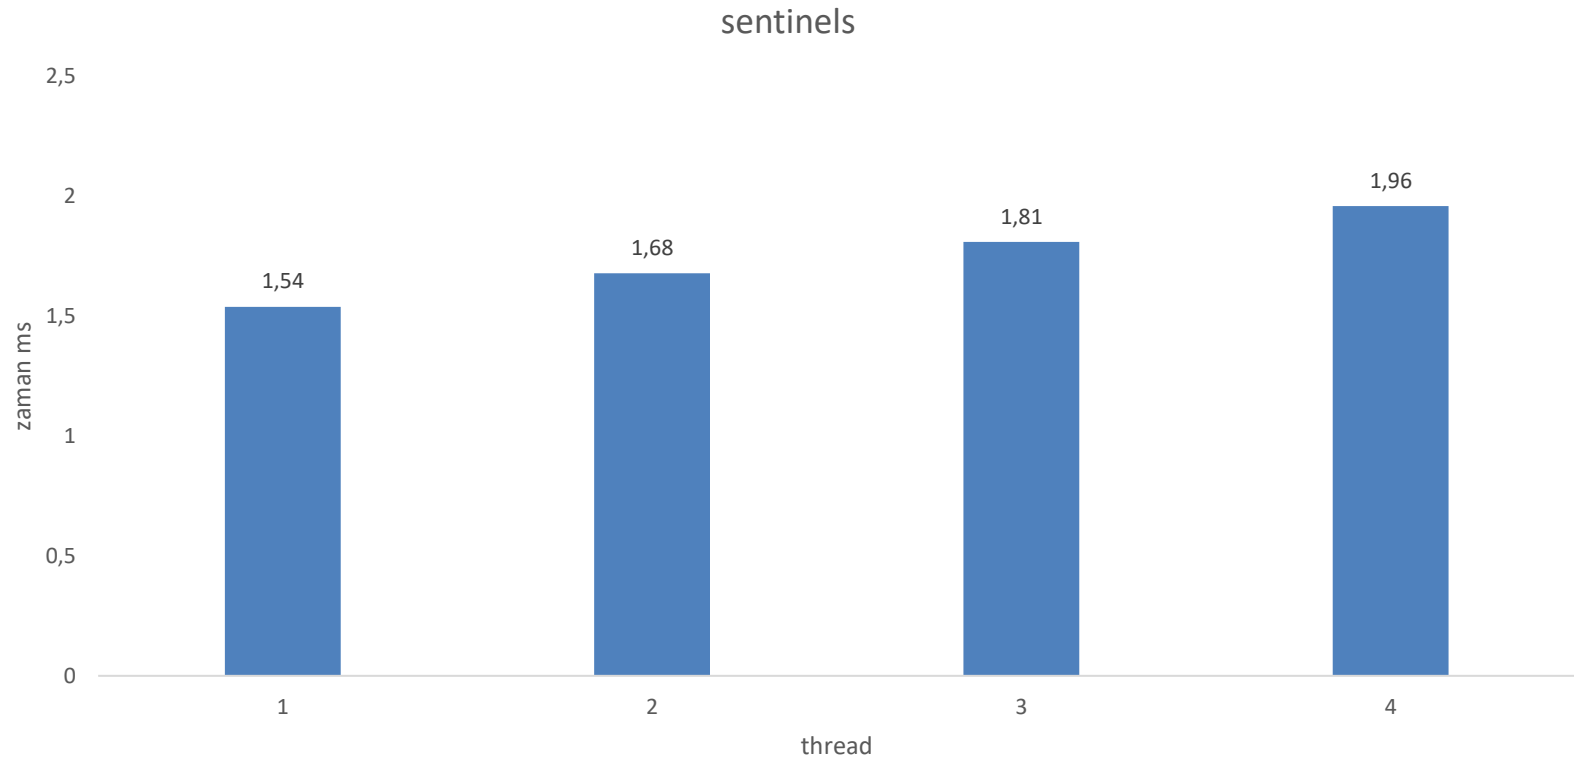
# Sonuç



## False Sharing – Sentinels

```
void work_with_sentinel(std::atomic<int>& a) {  
    thread_local unsigned int sentinel = 0;  
    for (int i = 0; i < 10,000,000; ++i)  
        ++sentinel;  
    a += sentinel;  
}  
  
int test_with_sentinels() {  
    std::atomic<int> a; a = 0;  
  
    std::thread t1([&] { work_with_sentinel(a); });  
    std::thread t2([&] { work_with_sentinel(a); });  
    std::thread t3([&] { work_with_sentinel(a); });  
    std::thread t4([&] { work_with_sentinel(a); });  
  
    t1.join(); t2.join(); t3.join(); t4.join();  
    return a;  
}
```

# Sonuç



**Teşekkürler**