

Synchronization And Communication



A Training for Software Development Professionals

Training Goals



- ☐ Understanding Real-Time Concepts
- ☐ Gaining the Knowledge of Real-Time Operating Systems and Internals

Contacts/Resources



- □Özgür Aytekin
 - ☐ Embedded Software Engineer at Collins Aerospace
 - □ozguraytekin@gmail.com

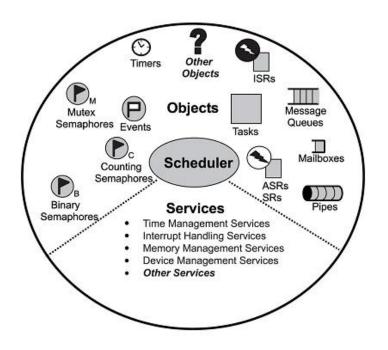


Chapter 1

Synchronization And Communication Kernel Objects

Agenda





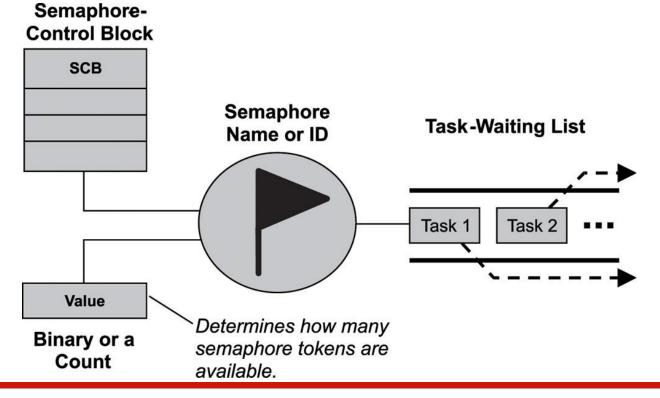
≻Semaphore

- ☐ Message Queues
- □ Pipes
- Connections
- Signals
- ■ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

Semaphore



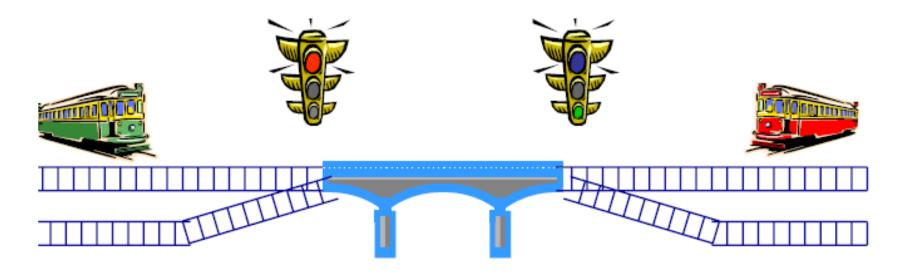
□ A semaphore (sometimes called a semaphore token) is a kernel object that one or more threads of execution can acquire or release for the purposes of synchronization or mutual exclusion.



Semaphore (cont)



☐ A "Semaphore" is a traffic light

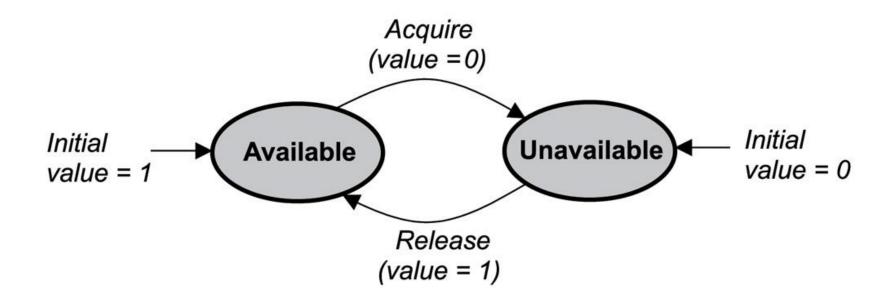


- ■Semaphores are a common solution in operating systems to deal with race conditions
- ☐ In real-time systems they lead to priority inversion

Binary Semaphores



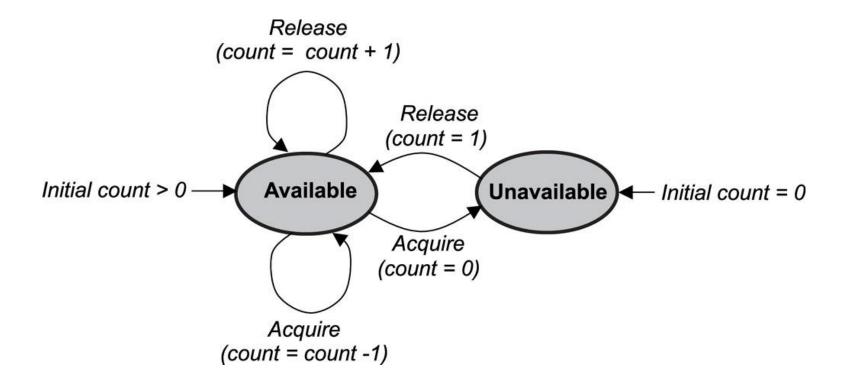
■A binary semaphore can have a value of either 0 or 1



Counting Semaphores



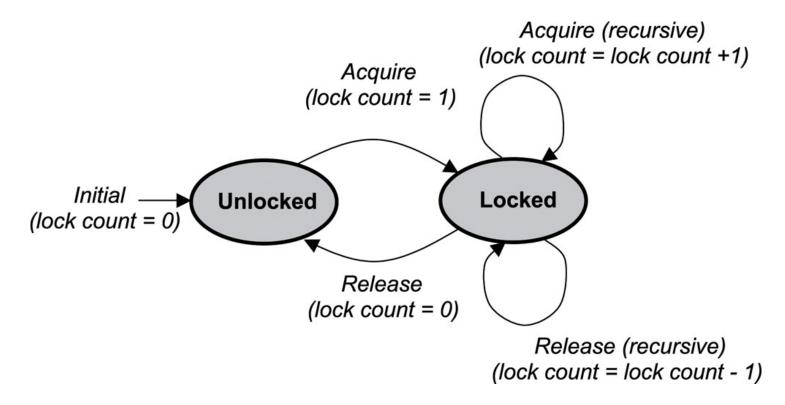
■A counting semaphore uses a count to allow it to be acquired or released multiple times



Mutual Exclusion (Mutex) Semaphores

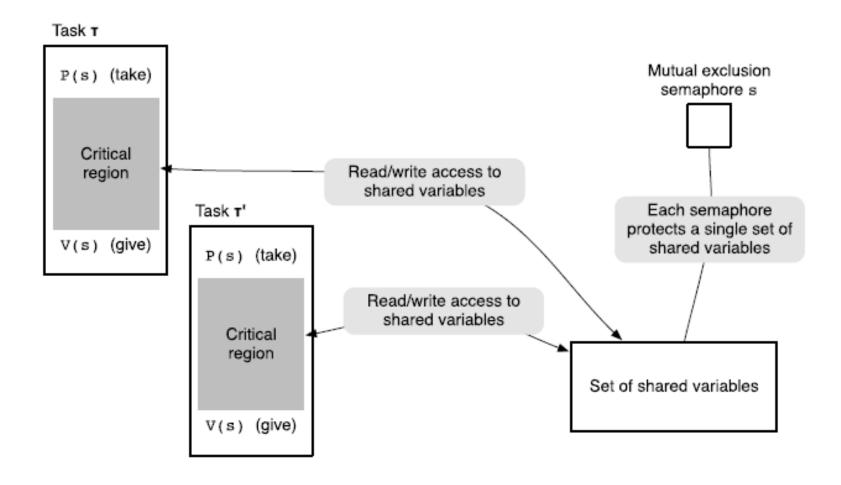


■A mutual exclusion (mutex) semaphore is a special binary semaphore that supports ownership, recursive access, task deletion safety



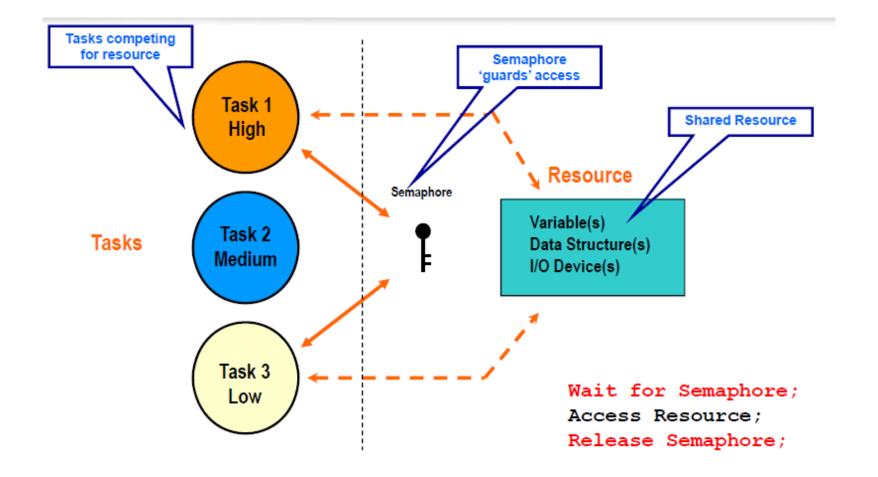
Usage of a semaphore for mutual exclusion





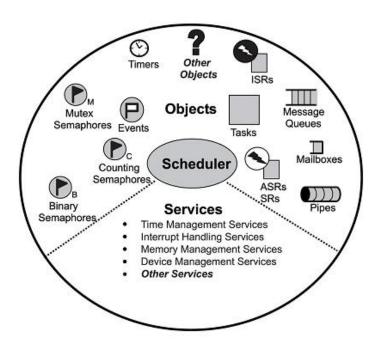
Usage of a semaphore for resource synchronization





Agenda



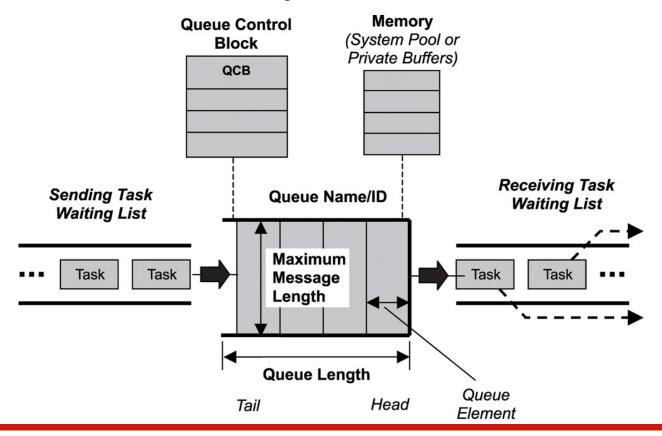


- ■Semaphore
- **≻Message Queues**
- Pipes
- Connections
- Signals
- ■ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

Message Queues

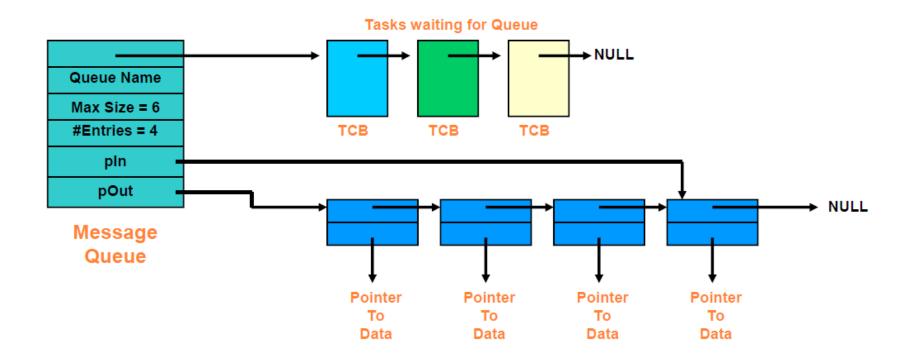


☐ A message queue is a buffer-like object through which tasks and ISRs send and receive messages to communicate and synchronize with data



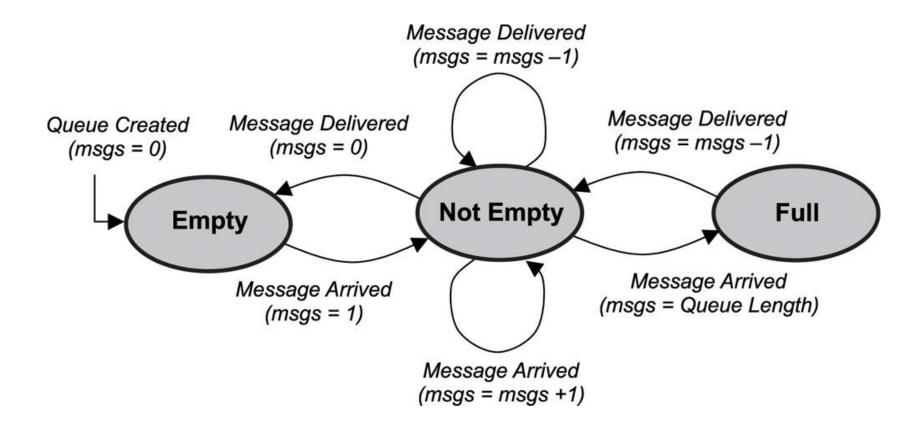
Message Queues - Implementation





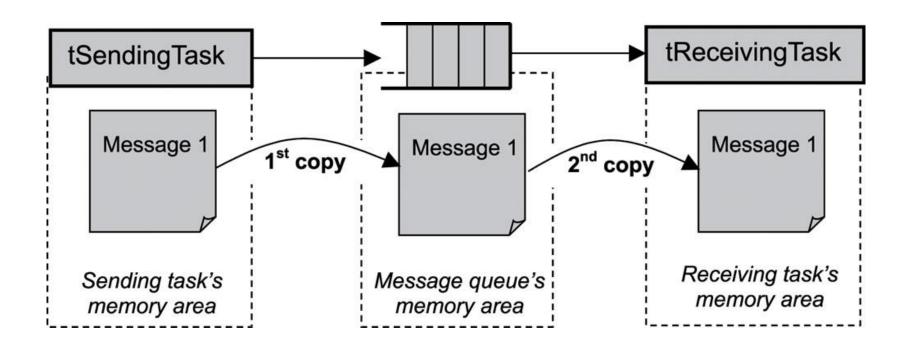
Message Queue States





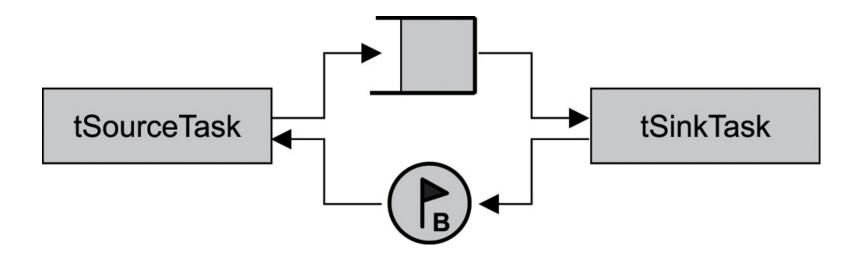
Message copying for sending and receiving messages





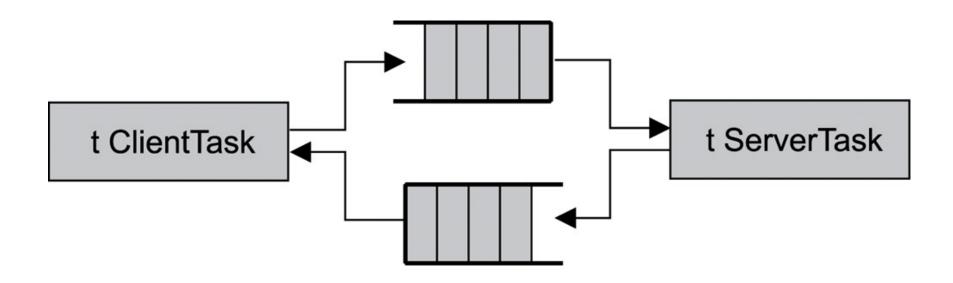
Non-Interlocked, One-Way Data Communication





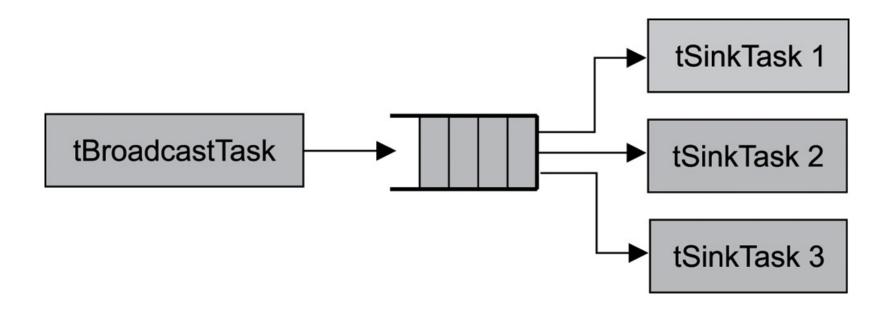
Interlocked, Two-Way Data Communication





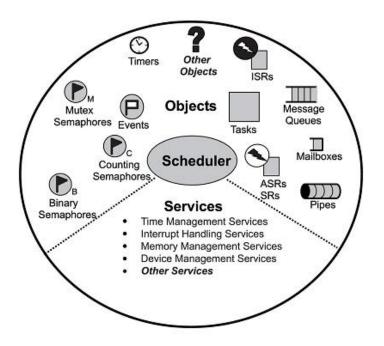
Broadcast Communication





Agenda



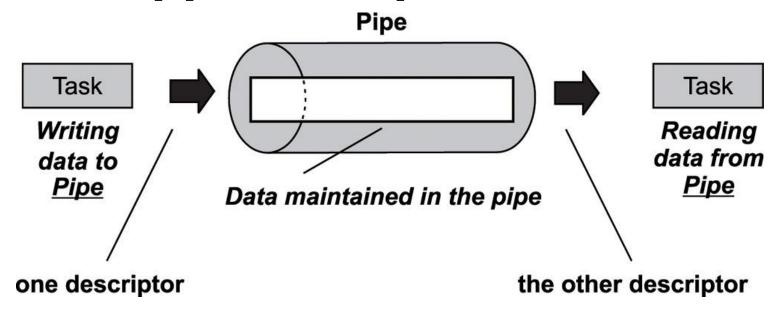


- ■Semaphore
- ☐ Message Queues
- **Pipes**
- ■Connections
- Signals
- □ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

Pipes




- ☐ Pipes are kernel objects that provide unstructured data exchange and facilitate synchronization among tasks
- ☐ Data consists of stream of bytes
- ☐ Data in a pipe cannot be prioritized




Common pipe operation





Pipe Control Blocks





Named and Unnamed Pipes

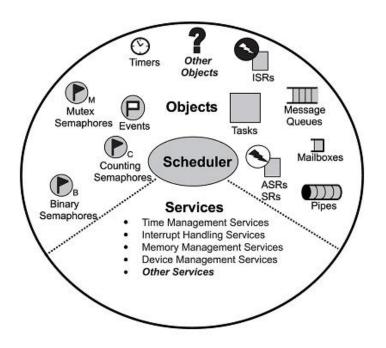


☐A named pipe has a name similar to a file name and appears in the file system as if it were a file or a device

- ☐ Any task or ISR that needs to use the named pipe can reference it by name
- ☐ The unnamed pipe does not have a name and does not appear in the file system

Agenda





- Semaphore
- ☐ Message Queues
- Pipes
- **Connections**
- **□**Signals
- ■ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

What is a Connection?



■ A Connection Object represents one end of a bidirectional communication channel

- ■A Connection Object represents a permission to communicate with whomever has the other Connection in the pair
- ☐ The other Connection in the pair can be in the same or a different Address Space

What is a Connection? (cont)

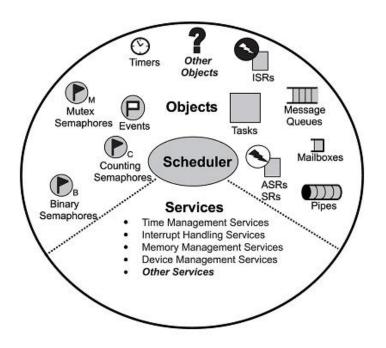


- ■A Connection is NOT a pipe
 - ■A pipe has a buffer –a Connection does not
 - A pipe is unidirectional –a Connection is bidirectional
- ☐ Data is copied directly from the sender to the receiver

☐ The Connection itself has no data storage

Agenda






- ■Semaphore
- ☐ Message Queues
- Pipes
- □ Connections
- **≻**Signals
- □ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

Signals



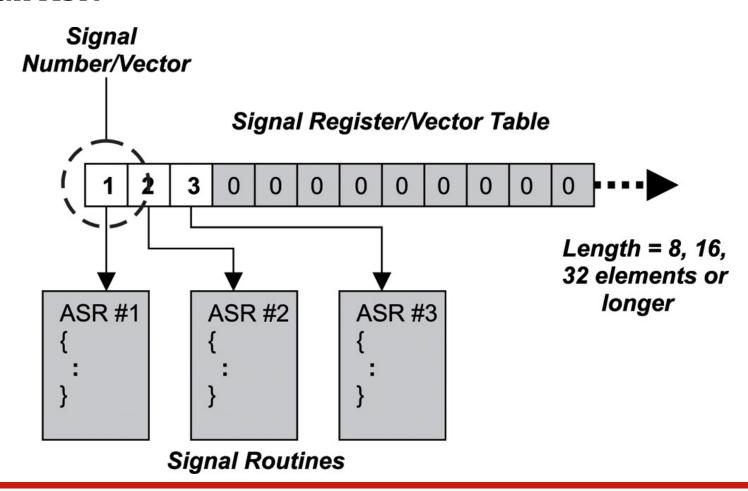
- A signal is a software interrupt that is generated when an event has occurred
- ☐ It diverts the signal receiver from its normal execution path and triggers the associated asynchronous processing.



Signal Vector Table

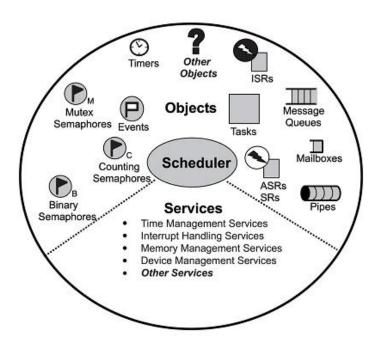


☐ Each element in the vector table is a pointer or offset to an ASR



Agenda





- ■Semaphore
- ☐ Message Queues
- Pipes
- Connections
- **□**Signals
- >ARINC 653 Ports
- ■Buffer and Blackboards
- ☐ Event Registers

APEX Channels



- ■A channel defines the following:
 - □Logical link between one source port and one or more destination ports
 - Mode of transfer of the messages from the source to the destination
 - ☐ Characteristics of the messages to be sent
- Each channel can be configured to operate in a specific mode
 - □sampling mode
 - queuing mode

Sampling Mode



Messages typically carry similar, but updated, data

- ■No queuing is performed
- ☐ A message remains in the source port until it is sent or overwritten.

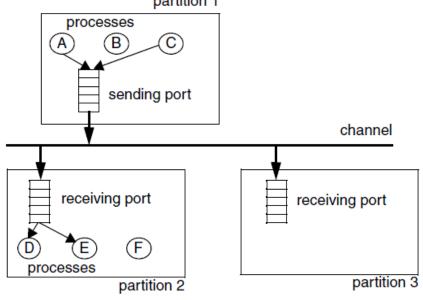
☐ The refresh rate indicates the maximum acceptable age of a valid message, from the time it was received at the port.

Queueing Mode



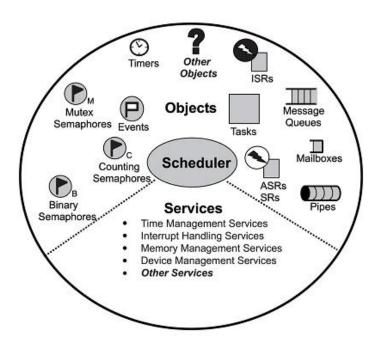
- ☐ Each new instance of a message may contain uniquely different data.
- Messages are queued in the source port until they are sent

☐ Messages are stored in the receiver port until a process reads them partition 1



Agenda





- Semaphore
- Message Queues
- Pipes
- □ Connections
- **□**Signals
- □ARINC 653 Ports
- **➤** Buffer and Blackboards
- ☐ Event Registers

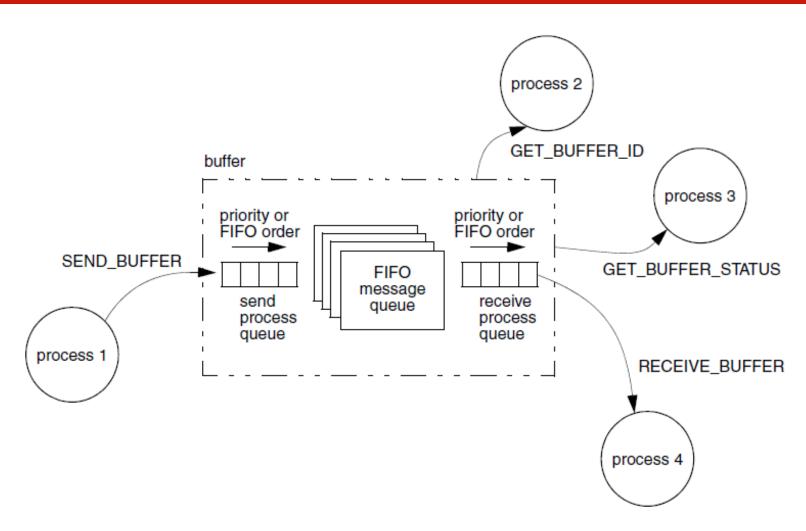
Buffers



- ☐Buffers store multiple messages in message queues and no messages are lost
- ■APEX buffers let processes communicate with each other within a partition
- ☐ A kind of message queue
- Asynchronous communication

Processes Using a Buffer to Communicate

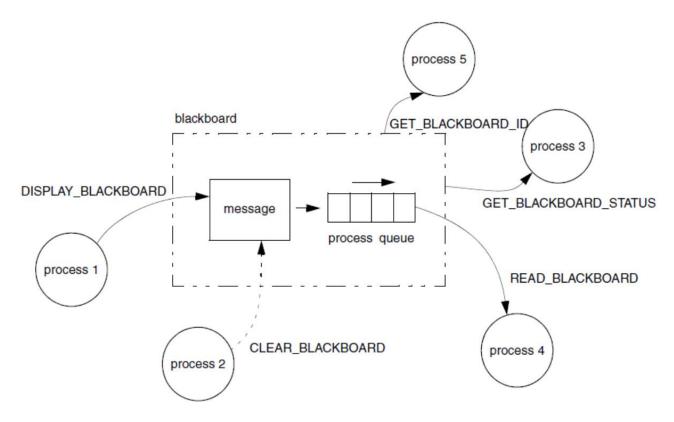




Blackboards

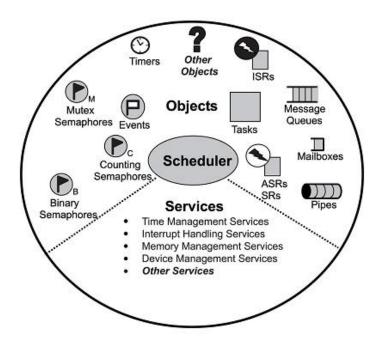


■Blackboards support a single message type between multiple source and destination processes



Agenda





- ■Semaphore
- ☐ Message Queues
- Pipes
- Connections
- **□**Signals
- ■ARINC 653 Ports
- ■Buffer and Blackboards
- **Event Registers**

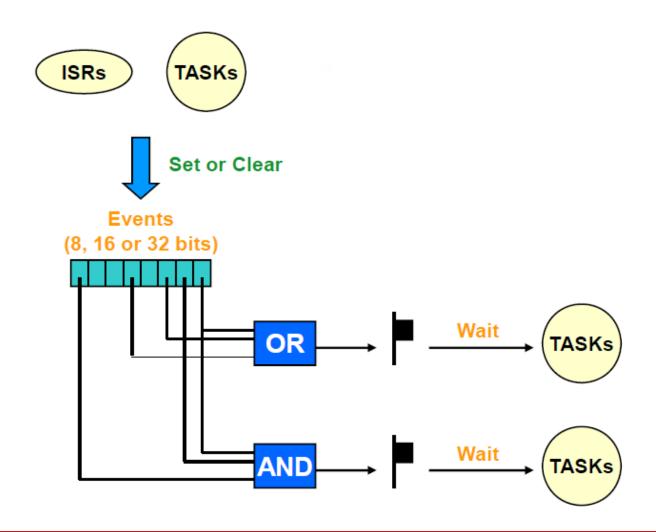
Event Registers



- ■Synchronization of tasks with the occurrence of multiple events
- ☐ Events are grouped
 - ■8, 16 or 32 bits per group
- ☐ Types of synchronization:
 - □ Disjunctive (OR): Any event occurred
 - □Conjunctive (AND): All events occurred
- ☐ Task(s) or ISR(s) can either Set or Clear event flags
- Only tasks can Wait for events

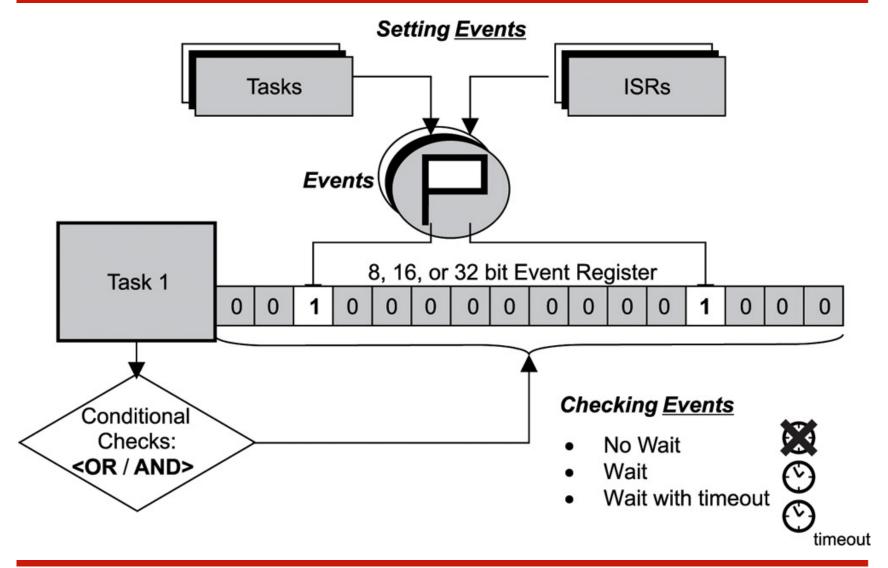
Event Registers





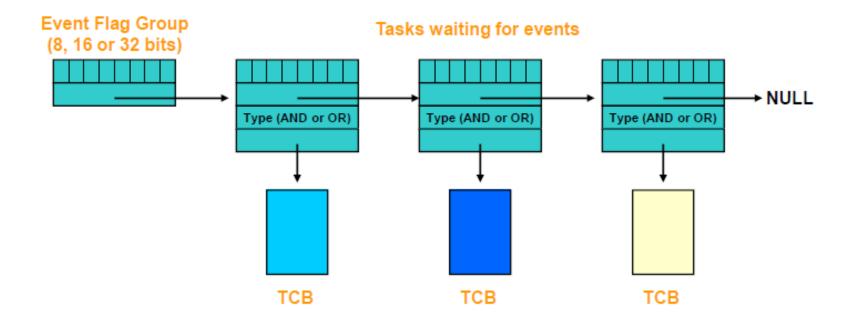
Event Registers





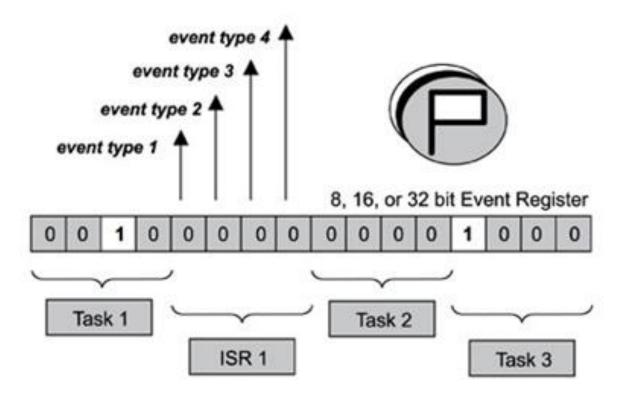
Event Registers - Implementation





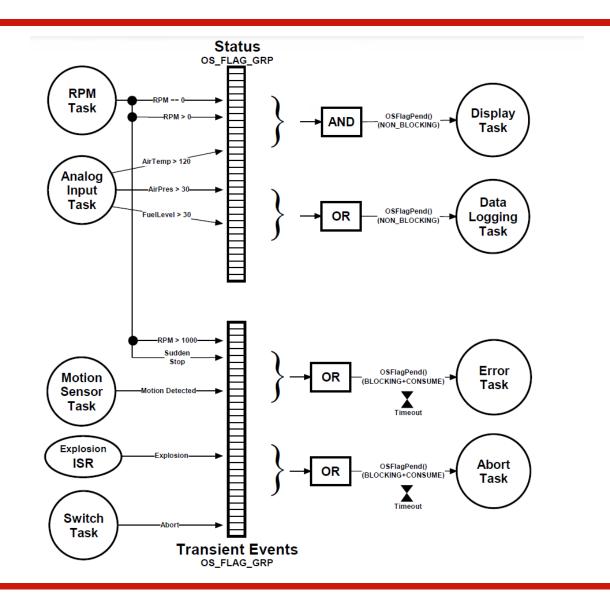
Identifying an event source





Event Registers - Example





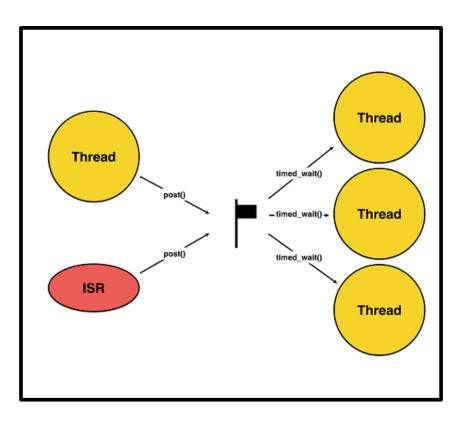


Chapter 2

Synchronization And Communication Methods

Agenda





>Synchronization

- **□**Communication
- ☐ Resource Synchronization
- ☐ Practical Design Patterns
- Deadlocks
- ☐ Priority Inversion

Synchronization



- ■Synchronization is classified into two categories:
 - ☐ Resource synchronization determines whether access to a shared resource is safe
 - □ Activity synchronization determines whether the execution of a multithreaded program has reached a certain state

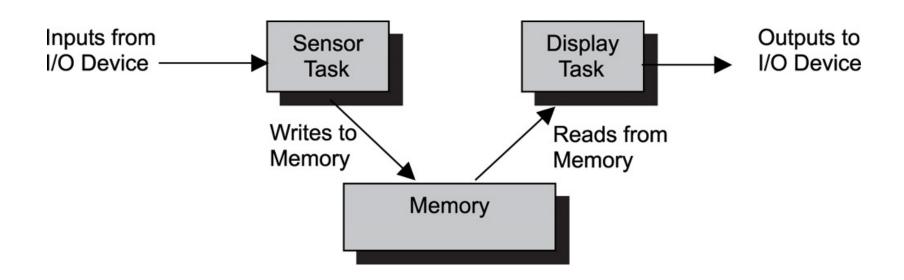


- ☐ Access by multiple tasks must be synchronized to maintain the integrity of a shared resource
- ☐ Mutual exclusion is a provision by which only one task at a time can access a shared resource

☐ A critical section is the section of code from which the shared resource is accessed.

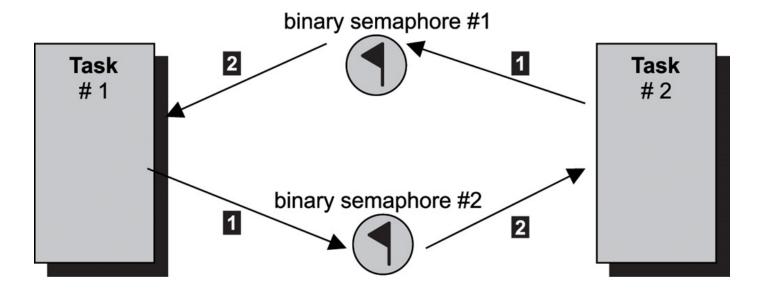


■ Multiple tasks accessing shared memory



Rendezvous synchronization





- ☐Both binary semaphores are initialized to 0
- ■When task #1 reaches the rendezvous, it gives semaphore #2, and then it gets on semaphore #1
- □When task #2 reaches the rendezvous, it gives semaphore #1, and then it gets on semaphore #2

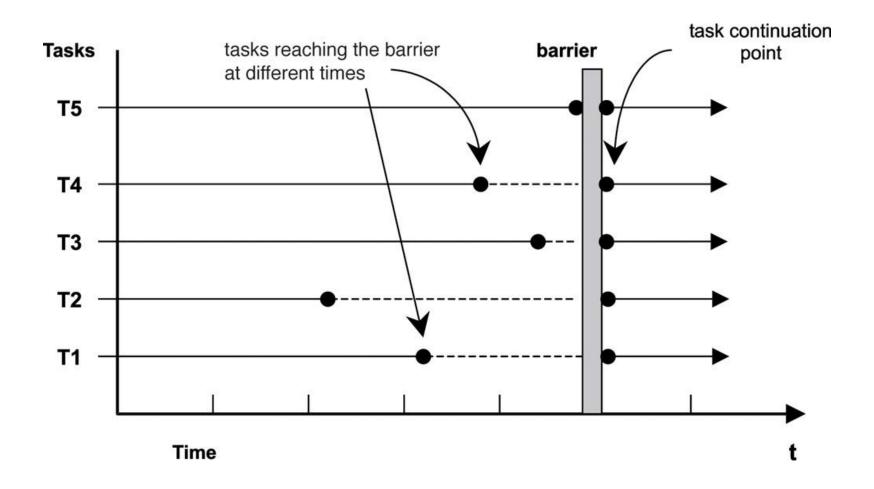
Activity Synchronization



- ■A task must synchronize its activity with other tasks to execute a multithreaded program properly
- □ Activity synchronization ensures that the correct execution order among cooperating tasks is used
- ☐ Activity synchronization can be either synchronous or asynchronous.

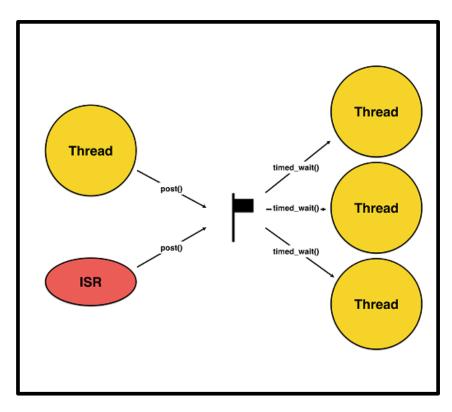
Barrier synchronization





Agenda





- Synchronization
- **Communication**
- ☐Resource Synchronization
- ☐ Practical Design Patterns
- □ Deadlocks
- ☐ Priority Inversion

Communication



- ☐ Tasks communicate with one another
 - ☐ To pass information
 - ☐ To coordinate their activities in a multithreaded embedded application
- □Communication can be signal-centric, datacentric, or both

Communication purposes

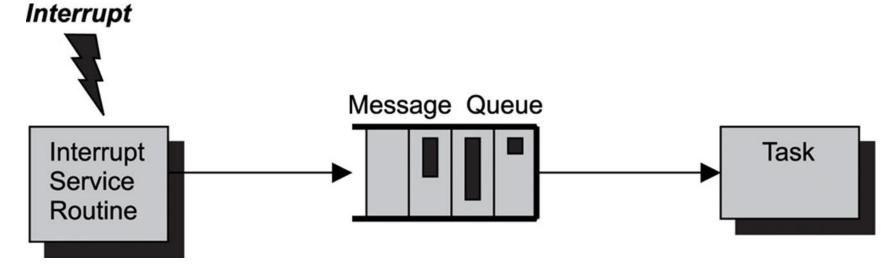


- ☐ Communication has several purposes, including the following:
 - ☐ Transferring data from one task to another,
 - ☐ Signaling the occurrences of events between tasks,
 - Allowing one task to control the execution of other tasks,
 - Synchronizing activities
 - ☐ Implementing custom synchronization protocols for resource sharing

Loosely coupled communication



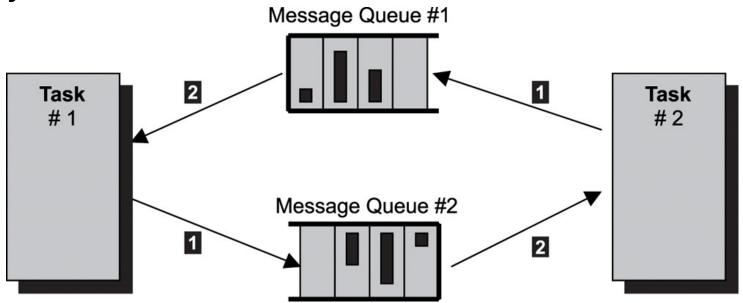
- ■When communication involves data flow and is unidirectional, this communication model is called loosely coupled communication
- Asynchronous



Tightly coupled communication

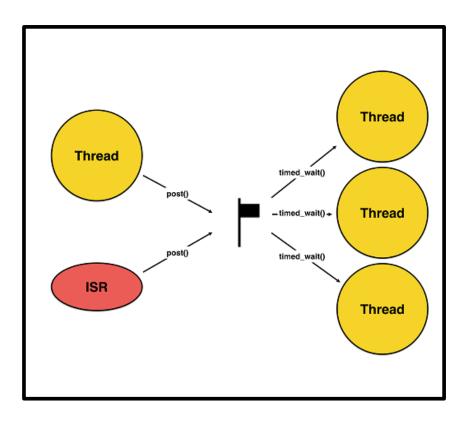


- ☐ Data movement is bidirectional
- ☐ The data producer synchronously waits for a response to its data transfer before resuming execution
- ■Synchronous



Agenda

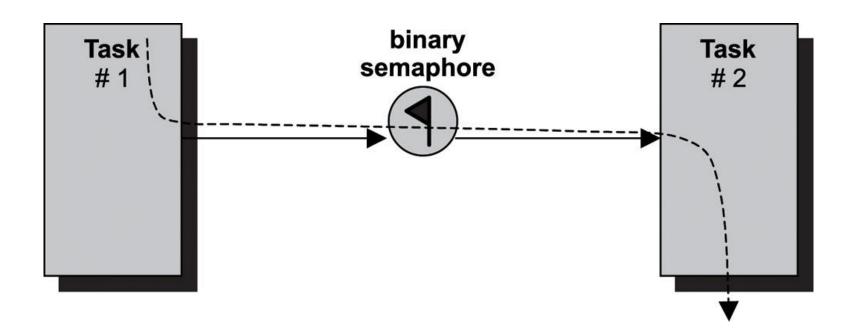




- Synchronization
- ■Communication
- ☐ Resource Synchronization
- ➤ Practical Design Patterns
- Deadlocks
- ☐Priority Inversion

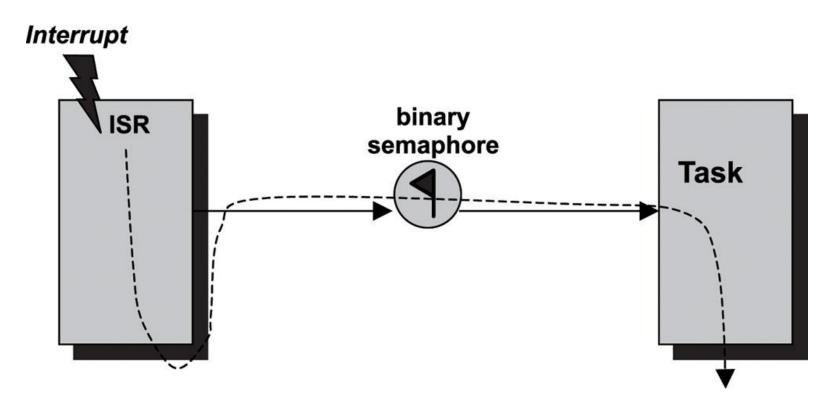


☐ Task-to-Task Synchronization Using Binary Semaphores






☐ ISR-to-Task Synchronization Using Binary Semaphores






☐ ISR-to-Task Synchronization Using Counting Semaphores



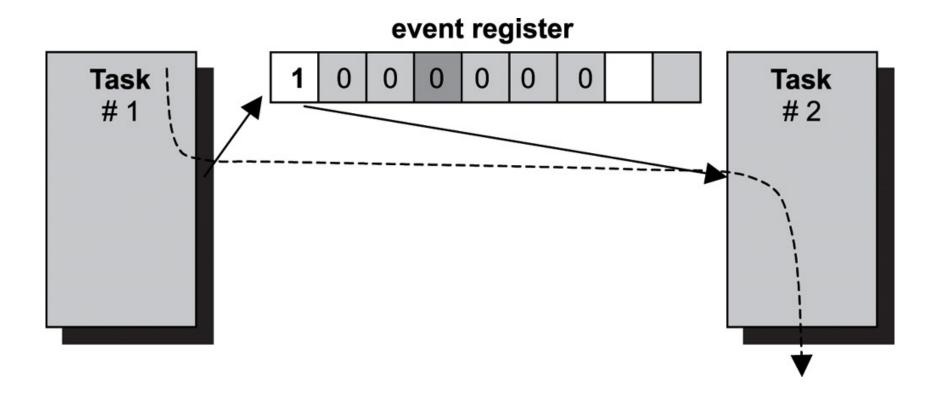


■Simple Rendezvous with Data Passing





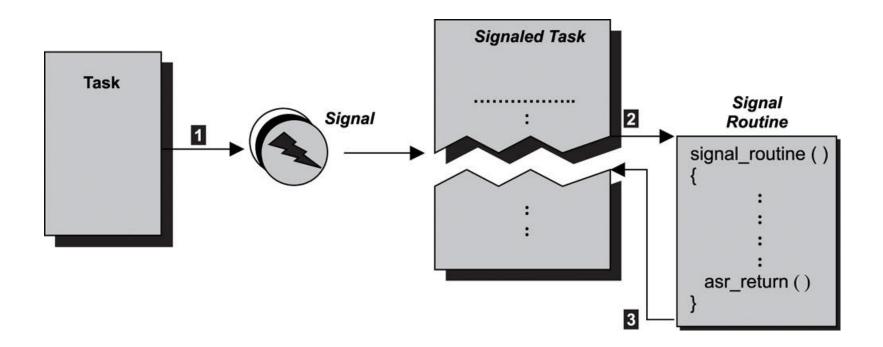
☐ Task-to-Task Synchronization Using Event Registers



Asynchronous Event Notification Using Signals



☐ Using signals for urgent data communication

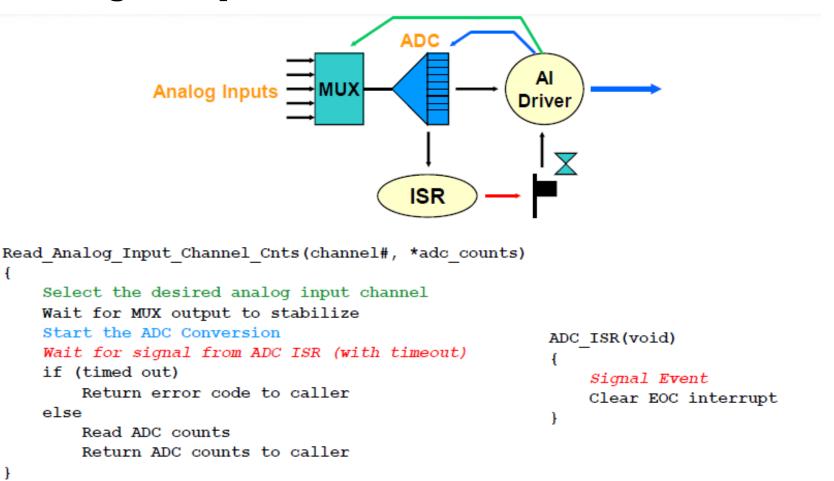


Asynchronous Event Notification Using Signals

else

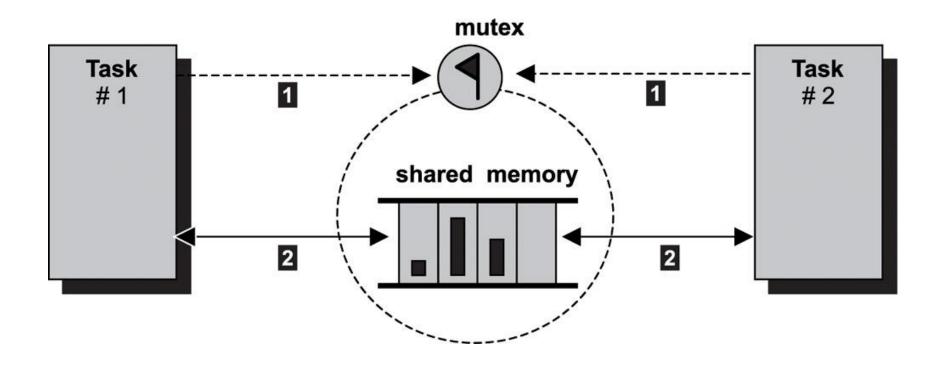


Using semaphores for data communication



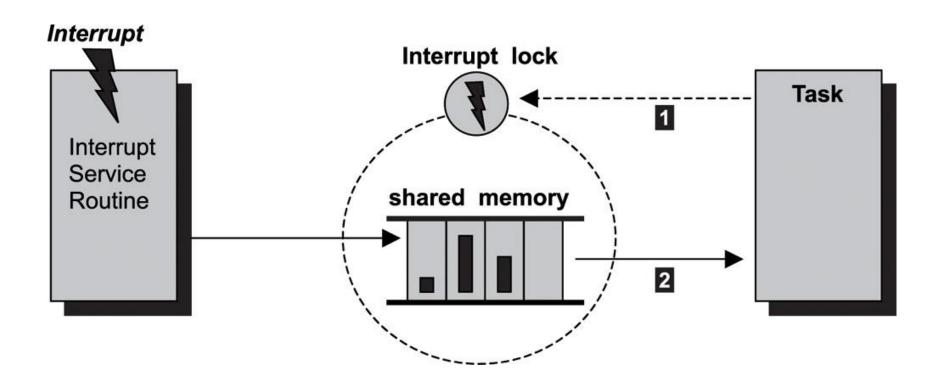


☐ Shared Memory with Mutexes






☐ Shared Memory with Interrupt Locks






☐ Shared Memory with Preemption Locks






☐ Sharing Multiple Instances of Resources Using Counting Semaphores and Mutexes





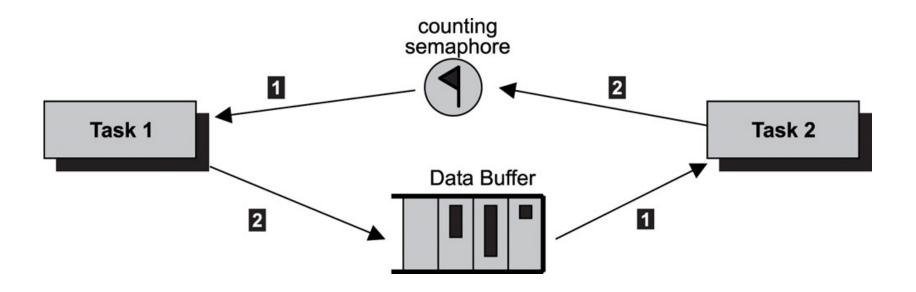
- Serial Communications (Rx)
- \square Rx ISR
 - ☐ Reads and buffers character
 - ☐ Clears interrupt source
 - □Signals semaphore (every character)



- □Rx Task
 - ■Wait on semaphore
 - □Gets character(s) from buffer

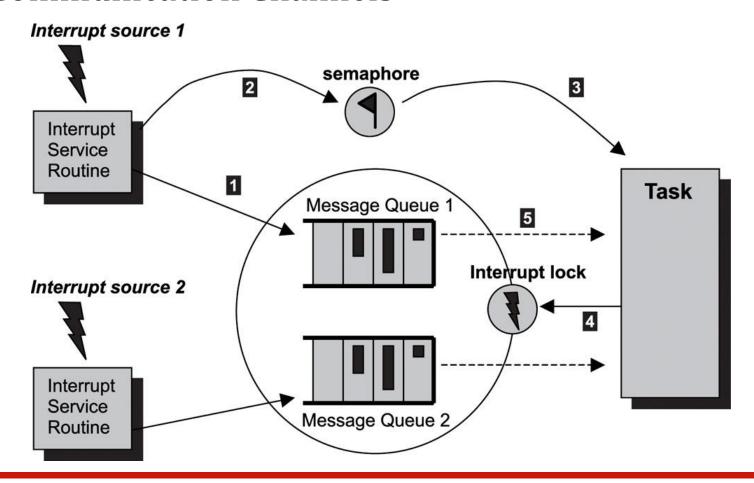


☐ Data Transfer with Flow Control






■ Asynchronous Data Reception from Multiple Data Communication Channels






☐ Sending High Priority Data between Tasks






☐ Message Queues & Buffer Pools






☐ Heap-Based Message Queue Communication between Tasks





- ☐ Error Handling with Message Queues
 - ☐ Error conditions are detected by tasks and ISRs then sent to an error handler
 - ☐ The error handler acts as a server



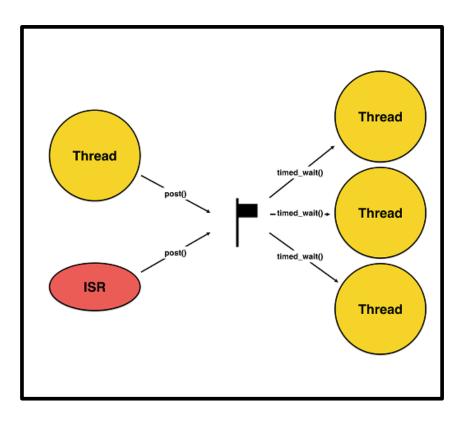


□I/O with Message Queue

```
Previous
                   Counts
                                      Message
                                                                         RPM
                             Delta
                                       Queue
                            Counts
                                                           RPM
                                                                         Avg. RPM
                    ISR
                                                           Task
                                                                        Under Speed
                                                  Counts
                      Counts=Fin * t
                                                                         Over Speed
                   32-Bit
                                                                         Max. RPM
         Fin
                  Input Capture
                                        RPMTask()
                                          while (1)
                                             Wait for message from ISR (with timeout);
RPM ISR()
                                            if (timed out)
                                                RPM = 0;
  Read Timer:
                                             else
  DeltaCounts = Counts
                                                RPM = 60 * Fin / counts;
               - PreviousCounts:
                                            Compute average RPM;
  PreviousCounts = Counts
                                            Check for overspeed/underspeed;
  Post DeltaCounts:
                                            Keep track of peak RPM;
                                             etc.
```

Agenda





- Synchronization
- ■Communication
- ☐Resource Synchronization
- ☐ Practical Design Patterns
- **Deadlocks**
- ☐Priority Inversion

Deadlocks



Deadlock is the situation in which multiple concurrent threads of execution in a system are blocked permanently because of resource requirements that can never be satisfied.

Deadlocks (cont)

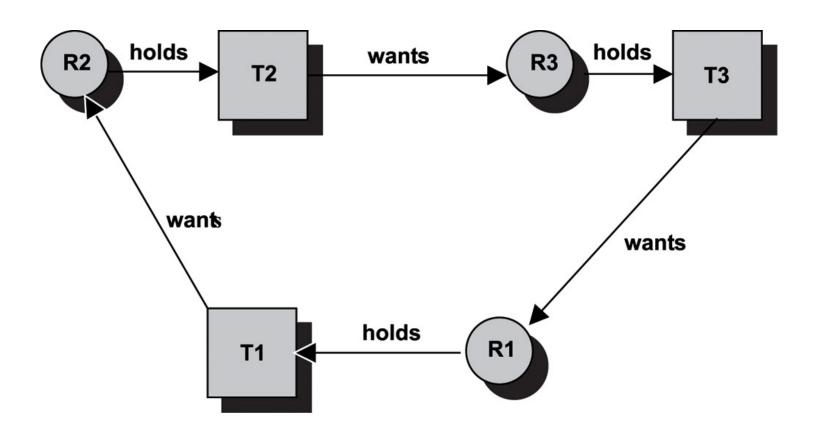


- ☐ Potential for deadlocks exist in a system in which the underlying RTOS permits resource sharing among multiple threads of execution
- □ Deadlock occurs when the following four conditions are present:
 - ■Mutual exclusion
 - ■No preemption
 - ☐ Hold and wait
 - ☐Circular wait

Deadlocks (cont)



☐ Deadlocks can involve more than two tasks.



A Deadlock example



```
void T1 (void)
   while (1) {
       Wait for event to occur;
                                      (1)
       Acquire M1;
                                        (2)
       Access R1;
                                        (3)
       \----- Interrupt!
                                      (4)
                                        (8)
       Acquire M2;
                                        (9)
       Access R2;
void T2 (void)
   while (1) {
       Wait for event to occur;
                                      (5)
      Acquire M2;
                                        (6)
       Access R2;
       Acquire M1;
                                        (7)
       Access R1;
```

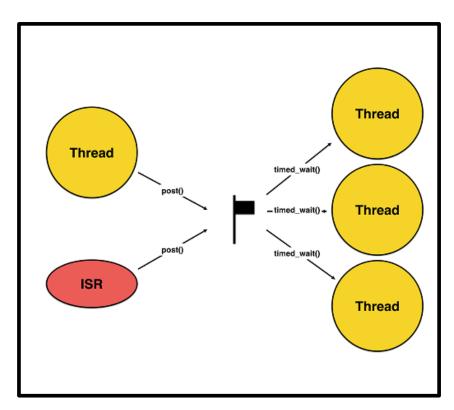
Some techniques used to avoid deadlocks



- ☐ Acquire all resources before proceeding
- ■Always acquire resources in the same order
- ☐ Use timeouts on wait calls

Agenda





- Synchronization
- **□**Communication
- ☐ Resource Synchronization
- ☐ Practical Design Patterns
- Deadlocks
- **▶**Priority Inversion

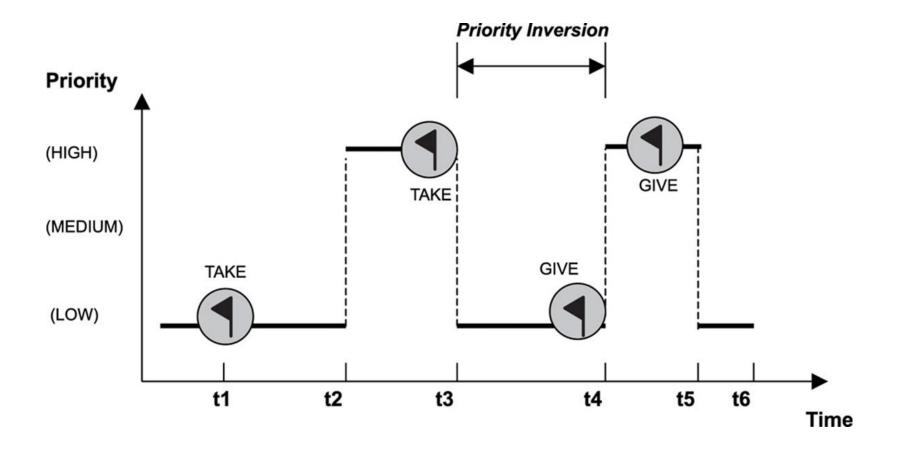
Priority Inversion



- □Blocking of a high priority task due to a lower priority task locking a shared resource.
- ☐ Priority inversion results from resource synchronization among tasks of differing priorities

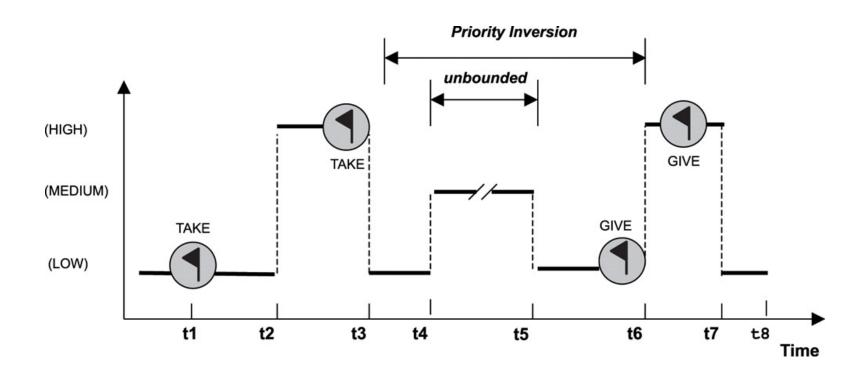
Bounded priority inversion





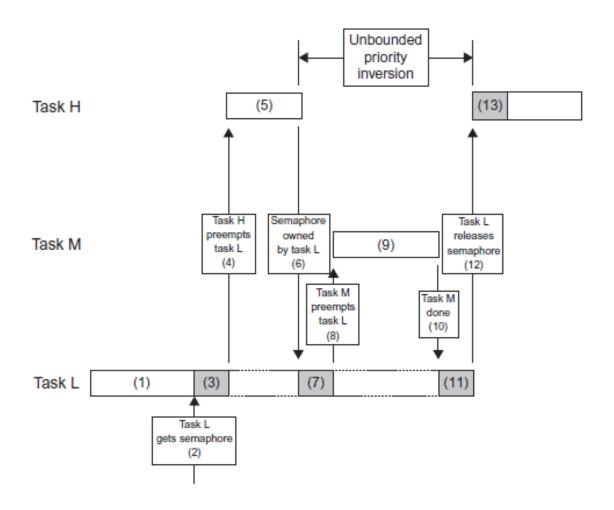
Unbounded priority inversion





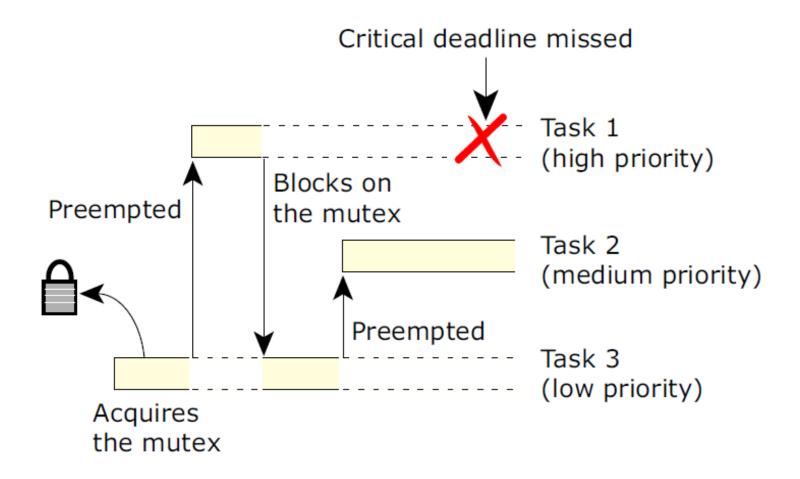
Unbounded priority inversion - 2





Another example of priority inversion





Solutions to Priority Inversion

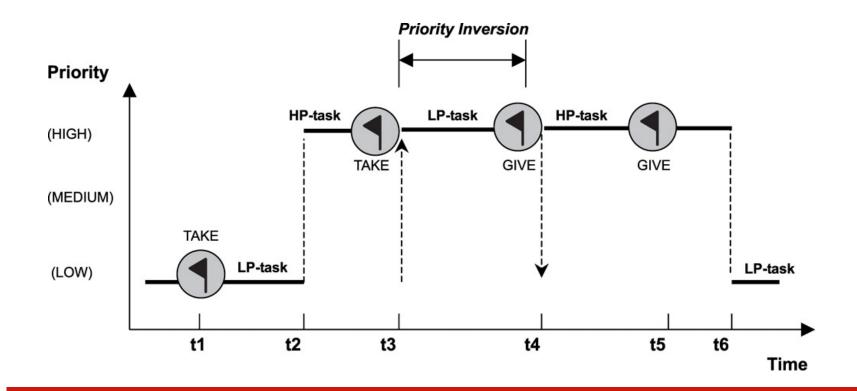


- ☐ There are two fundamental ways to deal with race conditions while avoiding priority inversion:
 - ☐ Use a mechanism that adjusts priorities of low priority tasks holding resources
 - ☐ Use a non-blocking mechanism

Priority Inheritance Protocol

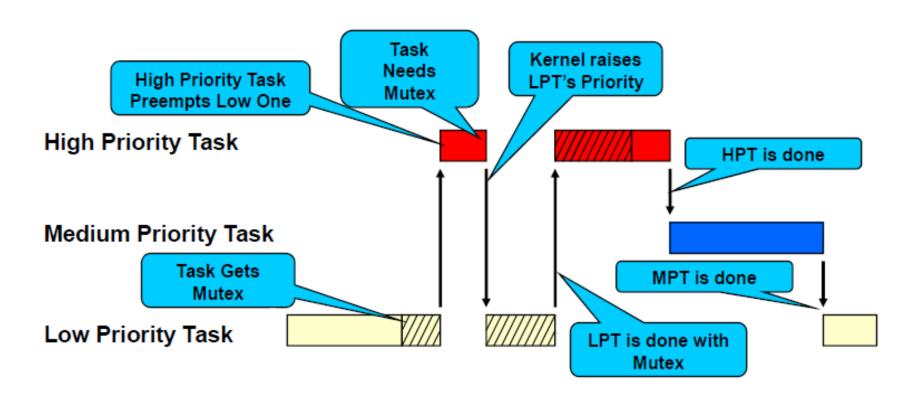


- ☐ A resource access control protocol that raises the priority of a task
- Not a good protocol



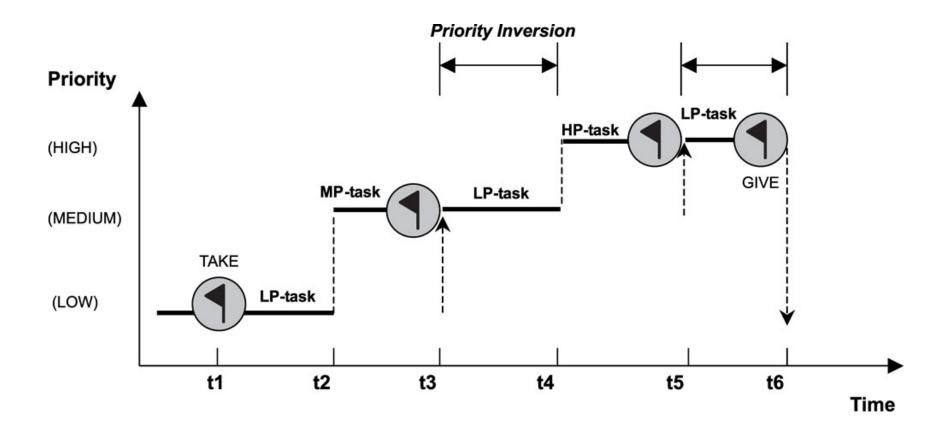
Mutex with Priority Inheritance





Transitive priority promotion

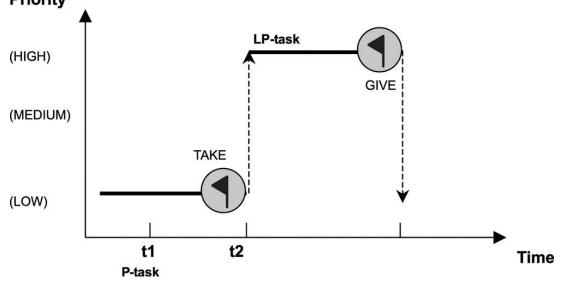




Highest Locking Protocol



- ☐ A fixed Priority, called priority ceiling is assigned to a resource
 - ☐ Priority is higher than any Task that contends for the resource
- ■When Task Obtains the resource
 - ☐ Task Priority is immediately elevated to priority ceiling



Priority Ceiling Protocol



- ☐ Highest Locking Protocol (HLP) is very similar
 - ☐ Similar to Priority Inheritance, except that each resource is given a priority ceiling, which is the priority of the highest-priority task that may lock the resource
 - ■When a resource is locked, task immediately inherits the priority ceiling
- ☐Advantage:
 - ☐It Works
- □ Disadvantages:
 - ☐ Significant RTOS overhead
 - □ Significant reduction in schedulable bound due to blocking priority ceiling